

# Many-task Computing with Python

**Monte Lunacek**  
**Jazek Braden**  
**Thomas Hauser**

Research Computing, University of Colorado **Boulder**

# What is many-task computing?

MTC, Raicu *et al.*

- set of tasks
  - may have dependencies
  - not necessarily single-core
- use a large resources over a short period

*High-throughput* Computing (HTC)

- executed over a period of months
- operations per month

# Why Python?

- flexible, powerful programming language
- easy to learn
- accessible to *non-programmers*
- glue
- free and open
- large community

# University of Colorado

- Some tightly-coupled MPI codes
- Many **independent tasks** (MTC)
- Diverse computing backgrounds
  - Geography
  - Ecology and Evolutionary Biology
  - Microbial Ecology
  - Astronomy
  - Geology
- Range of computational experience

# Supercomputing Without the Pain

- Accessible to anyone with:
  - Simulation or analysis to run
  - Desire to do it faster
- Remove barriers to entry

# Outline

MTC

Solutions, **mpi4py** example

Message queues

- **IPython**
- **Celery**
- Fault-tolerance, elasticity, memory

Scaling Results

Conclusions

# Many-task Computing

```
def work(x):  
    import time, os  
    start_time = time.time()  
    time.sleep(x)  
    end_time = time.time()  
    return {'pid': os.getpid(),  
            'start_time': start_time,  
            'end_time': end_time}
```

Monte Carlo simulations

Parameter scan

Uncertainty Quantification

Parameter Optimization

# Approach

Condor/Moab/SLURM

Bash/pbsdsh

MPI **mpi4py**

Message broker: **IPython, Celery**

SAGA BigJob

and many, many more ...




# Bash/pbsdsh

```
#!/bin/bash  
PATH=$PBS_O_WORKDIR:$PBS_O_PATH  
TRIAL=$(( $PBS_VNODENUM + $1 ))  
python work.py 5
```

```
for i in {1..N}  
do  
    pbsdsh wrapper.sh $count  
    count=$(( $count + 12 ))  
done
```

A little **painful**

# A little inefficient




# MPI example





# Message queues



fault-tolerance, elasticity, memory



# Weak Scaling

Compare: mpi4py, IPython Parallel, Celery

What recommendations can we offer?

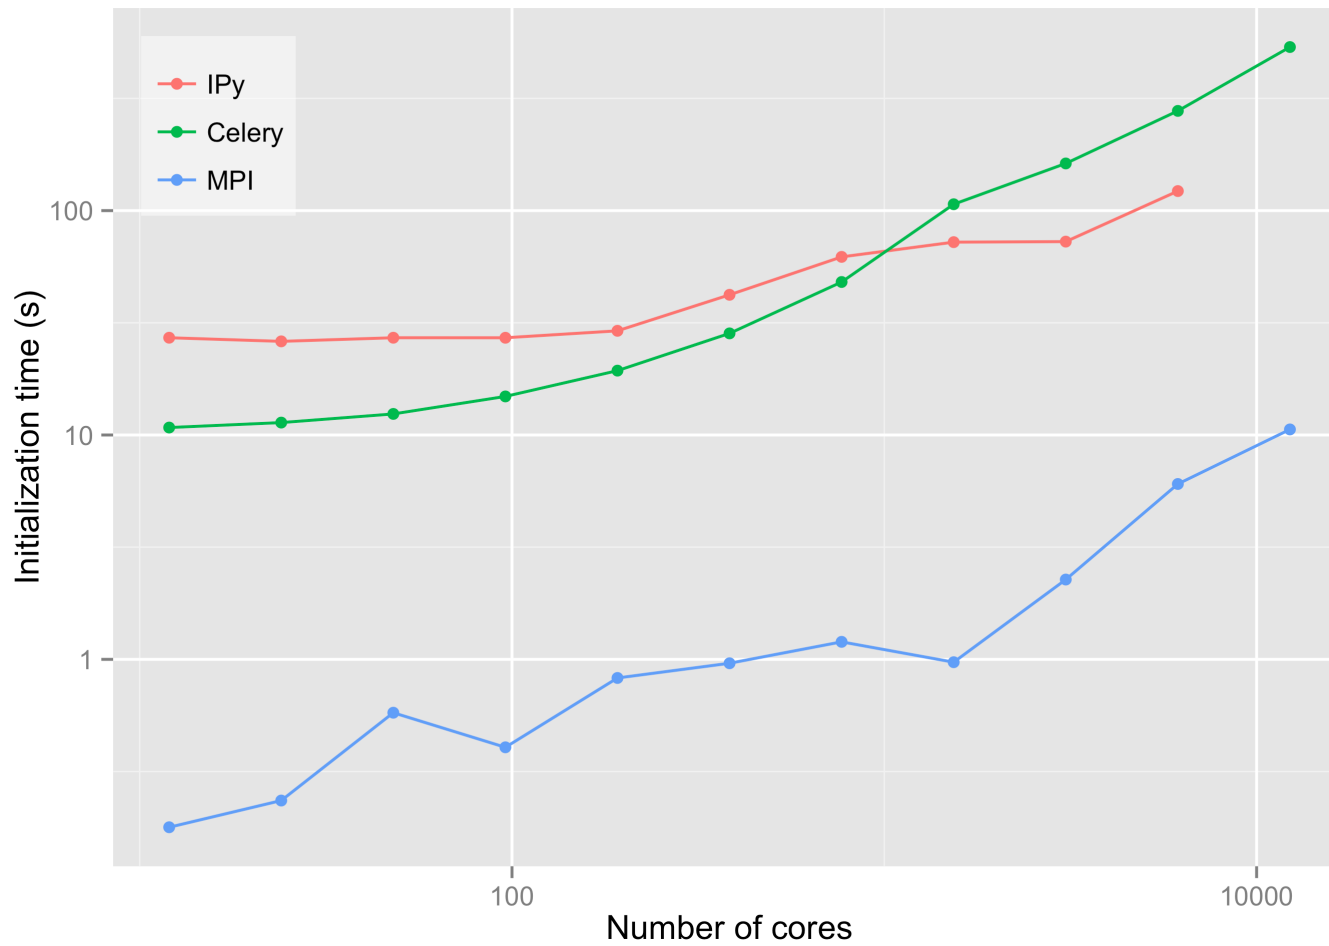
Best case scenario

- small messages
- limited file IO
- sleep

Weak scaling


- $\text{task} = 10 * \text{number of cores}$
- up to 12,288 cores
- time uniform(27,33) seconds

# Initialization

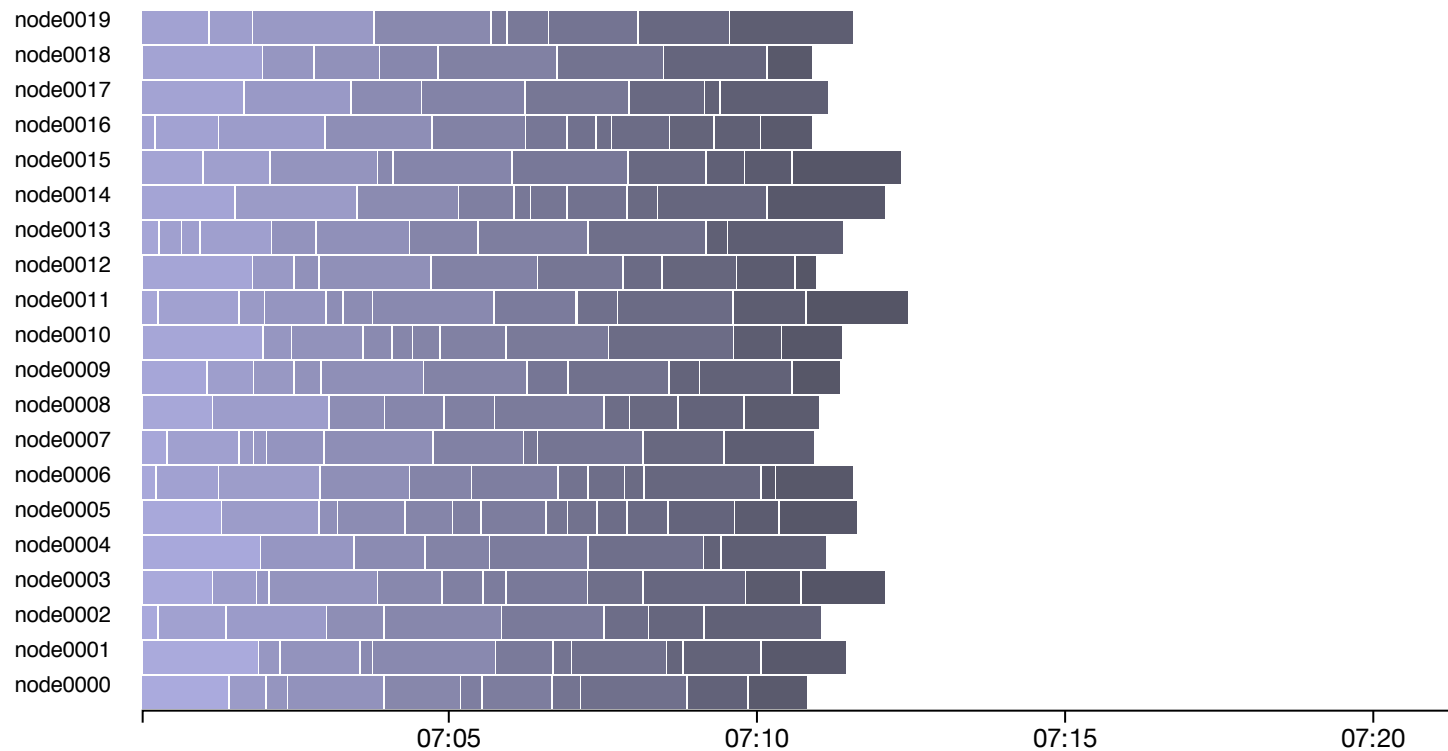




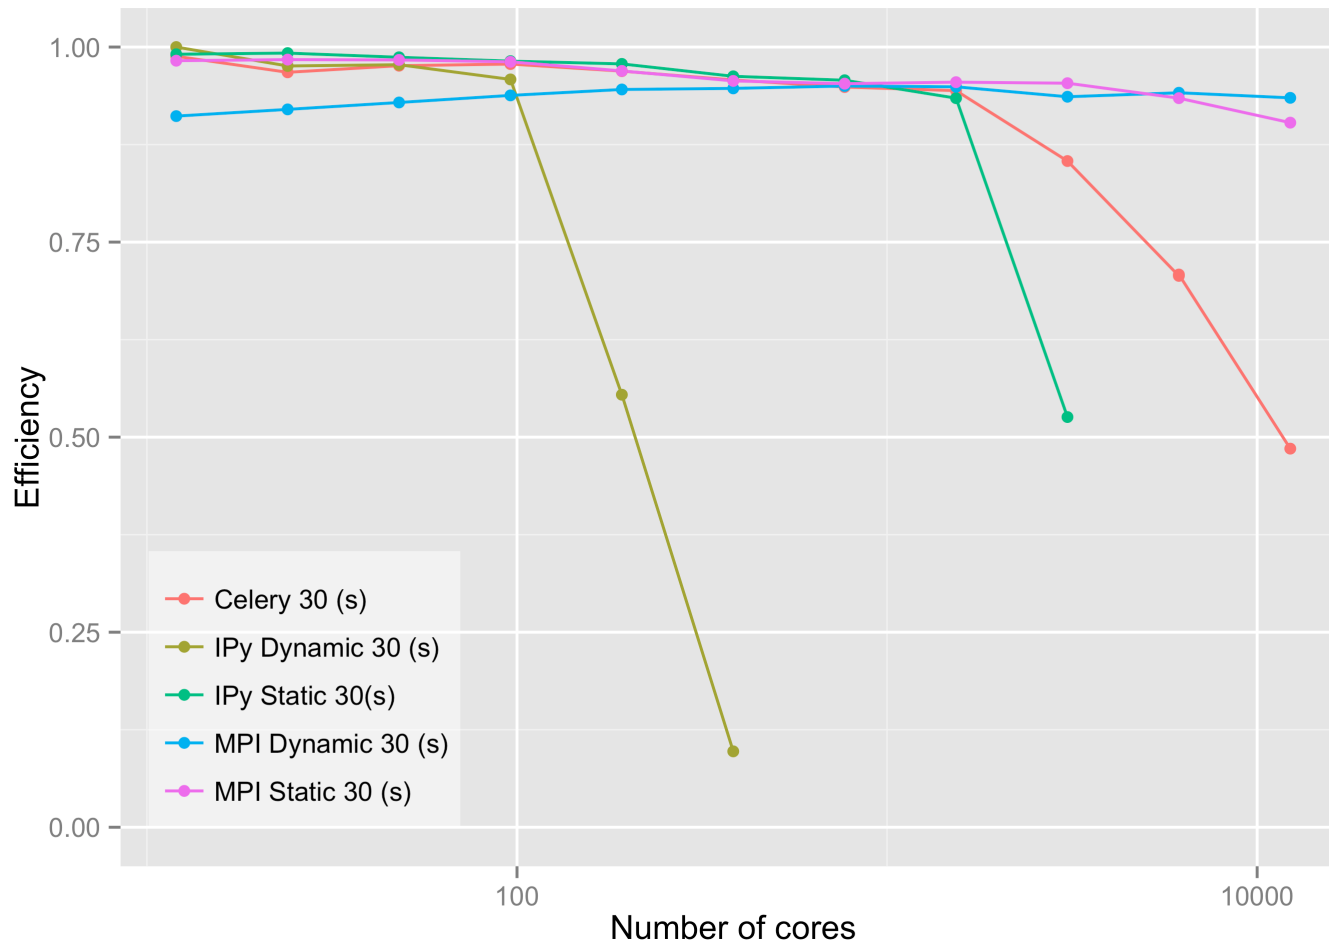
# Scheduling: Static



# Scheduling: Dynamic



# Weak Scaling



# Compare

## Recommend

mpi4py	many cores, robust simulation
IPython	100 cores, not robust
IPython	many cores, not robust, consistent time
Celery	many cores, not robust, variable time
Multiprocessing	single-node

## Issues

IPython, Celery	launching
all	user interface

[https://github.com/mlunacek/load\\_balance\\_ipython](https://github.com/mlunacek/load_balance_ipython)

([https://github.com/mlunacek/load\\_balance\\_ipython](https://github.com/mlunacek/load_balance_ipython))

# Conclusions

Python is an excellent way to manage MTC jobs

IPython and Celery

- elastic
- memory
- **fault-tolerant**

Moving forward

- Simple interface for MTC
- Profiling tools
- Realistic tasks

# References

## Paper

*Scaling of Many-Task Computing Approaches in Python on Cluster Supercomputers*

Monte Lunacek *et al.* IEEE Cluster 2013

## Tutorial

tomorrow: 8:30 - 10:15, room 9

## slides

[https://github.com/mlunacek/python\\_frcrc\\_2013](https://github.com/mlunacek/python_frcrc_2013)

## links

Distributions: Enthought

(<http://www.enthought.com/products/epd.php>), Anaconda

(<https://store.continuum.io/cshop/anaconda>)

Packages: IPyθοParallel (<http://ipython.org/ipython->